





```
In [ ]: from sklearn.feature_selection import SelectKBest, chi2
from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()
X_scaled = scaler.fit(X).transform(X)
X_scaled

Out [ ]: array([[0.89073462, 0.80987235, 0.8779829 , ..., 0.72740866, 0.82026832,
        0.85913693],
       [0.10258527, 0.78723745, 0.21941259, ..., 0.24100718, 0.04457839,
        0.07310273],
       [0.09913998, 0.71626466, 0.2556253 , ..., 0.26721914, 0.07971828,
        0.06099862 ],
       [0.86894928, 0.75106377, 0.86598024, ..., 0.72631794, 0.8528543 ,
        0.81182299],
       [0.04070846, 0.46993384, 0.05424333, ..., 0.24168541, 0.27271645,
        0.34995666],
       [0.94234337, 0.87639288, 0.86258877, ..., 0.61200446, 0.75515142,
        0.79526219]])

In [ ]: print(X_scaled[0]) # for reference

0.89073462 0.80987235 0.8779829 0.71654357 0.87368383 0.78048193
0.8518387 0.87368383 0.78048193 0.83263085 0.87148368 0.6727923
0.72740866 0.82026832 0.85913693

In [ ]: from sklearn.feature_selection import SelectKBest, chi2
X_new = SelectKBest(chi2, k=5).fit_transform(X_scaled, y)
X_new

Out [ ]: array([[0.89073462, 0.8779829 , 0.71654357, 0.87368383, 0.78048193],
       [0.10258527, 0.21941259, 0.93923032, 0.03176728, 0.07591281],
       [0.09913998, 0.2556253 , 0.92452282, 0.06231308, 0.08937904],
       ...,
       [0.86894928, 0.86598024, 0.73817325, 0.88009701, 0.76335731],
       [0.04370846, 0.05424333, 0.08953574, 0.11858609, 0.3650229],
       [0.94234337, 0.86258877, 0.71981219, 0.86702489, 0.79852367]])

In [ ]: print(X_new[0])

[0.89073462 0.8779829 0.71654357 0.87368383 0.78048193]
```

If we compare the above result with the first row of the scaled data, we can map the column names as : ACDHI

Therefore, based on the selecting features on distance similarity, we have come up with the 5 best features to choose for our model : [A, C, D, H, I]. Let us try to build our pipeline using these new features only.

```
In [ ]: new_df = df.copy(['A', 'C', 'D', 'H', 'I', 'Class'])
new_df.head()
```

```
Out [ ]:
   A      C      D      H      I  Class
0 231.420023 217.624839 -15.611916 198.160805 82.873279 2
1 -38.019270 9.583547 22.293822 -33.711852 -8.356041 1
2 -39.197085 -20.418850 21.023083 19.790280 -25.992919 -6.612401 3
3 221.630408 216.723322 -9.900781 197.640135 82.560019 2
4 228.558412 204.637218 -13.277704 209.300011 89.961688 3
```

```
In [ ]: X_new = new_df.drop(['Class'], axis=1)
y_new = new_df['Class']
X_new_train, X_new_test, y_new_train, y_new_test = train_test_split(X_new, y_new, stratify=y, test_size=0.2)
```

```
In [ ]: # Decision Tree Pipeline
pipeline_dt_new = Pipeline([('scaler', StandardScaler()),
                             ('pca', PCA(n_components=2)),
                             ('dt_classifier', DecisionTreeClassifier(max_depth=3))])
pipeline_dt_new.fit(X_new_train, y_new_train)
pipeline_dt_new.score(X_new_test, y_new_test)
```

```
Out [ ]: 0.4993206333333333
```

Unfortunately, we do not notice any improvements in accuracy using the K-Best distance similarity metric to reduce the number of features to the most important features.

My next approach was to try different model pipelines on the data and see if I can receive a better accuracy. The models I tried were :

- Logistic Regression
- Decision Tree Classifier
- Random Forest Classifier
- Naive-Bayes Classifier

```
In [ ]: ## Creating pipelines for the different classifiers
## Logistic Regression
pipeline_lr = Pipeline([('scaler1', StandardScaler()),
                        ('pca1', PCA(n_components=2)),
                        ('lr_classifier', LogisticRegression(random_state=10))])
pipeline_lr.fit(X_train, y_train)
score = pipeline_lr.score(X_test, y_test)
pred = pipeline_lr.predict(X_test)
acc1 = metrics.accuracy_score(y_test, pred)
print('Classification report:')
print(metrics.classification_report(y_test, pred))
print('Accuracy score: ', acc)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/validation.py:760: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), f
or example using ravel().
  y = column_or_1d(y, warn=True)
Score: 0.4991777777777778
Classification report:
precision    recall  f1-score   support
1         0.00         0.00         0.00         59804
2         0.50         1.00         0.67         179704
3         0.00         0.00         0.00         120492
accuracy          0.25          0.33          0.33         360000
macro avg         0.17         0.33         0.22         360000
weighted avg         0.25         0.50         0.33         360000
Accuracy: 0.4991777777777778
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarni
ng: Precision and f-score are ill-defined and being set to 0.0 in labels with no predicted samples. Us
e 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
precision    recall  f1-score   support
1         0.00         0.00         0.00         59804
2         0.50         1.00         0.67         179704
3         0.00         0.00         0.00         120492
accuracy          0.25          0.33          0.33         360000
macro avg         0.17         0.33         0.22         360000
weighted avg         0.25         0.50         0.33         360000
Accuracy score: 0.4991722222222224
```

```
In [ ]: ## Decision Tree
pipeline_dt = Pipeline([('scaler2', StandardScaler()),
                        ('pca2', PCA(n_components=2)),
                        ('dt_classifier', DecisionTreeClassifier(max_depth=3, random_state=0))])
pipeline_dt.fit(X_train, y_train)
score2 = pipeline_dt.score(X_test, y_test)
pred2 = pipeline_dt.predict(X_test)
acc2 = metrics.accuracy_score(y_test, pred2)
print('Classification report:')
print(metrics.classification_report(y_test, pred2))
print('Accuracy score: ', acc2)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1272: UndefinedMetricWarni
ng: Precision and f-score are ill-defined and being set to 0.0 in labels with no predicted samples. Us
e 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
precision    recall  f1-score   support
1         0.00         0.00         0.00         59804
2         0.50         1.00         0.67         179704
3         0.25         0.00         0.00         120492
accuracy          0.33          0.33          0.33         360000
macro avg         0.25         0.33         0.22         360000
weighted avg         0.25         0.50         0.33         360000
Accuracy score: 0.4991722222222224
```

```
In [ ]: ## Random Forest
pipeline_rf = Pipeline([('scaler3', StandardScaler()),
                        ('pca3', PCA(n_components=2)),
                        ('rf_classifier', RandomForestClassifier(max_depth=5, n_estimators=5))])
pipeline_rf.fit(X_train, y_train)
score3 = pipeline_rf.score(X_test, y_test)
pred3 = pipeline_rf.predict(X_test)
acc3 = metrics.accuracy_score(y_test, pred3)
print('Classification report:')
print(metrics.classification_report(y_test, pred3))
print('Accuracy score: ', acc3)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/pipeline.py:354: DataConversionWarning: A column-vecto
r y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for exampl
e using ravel().
  self._final_estimator.fit(Xt, y, **fit_params)
Classification report:
precision    recall  f1-score   support
1         0.12         0.00         0.00         59804
2         0.50         1.00         0.67         179704
3         0.33         0.00         0.00         120492
accuracy          0.32         0.33         0.32         360000
macro avg         0.38         0.50         0.33         360000
Accuracy score: 0.4991388888888889
```

```
In [ ]: ## Naive Bayes Classifier
pipeline_nb = Pipeline([('scaler4', StandardScaler()),
                        ('pca4', PCA(n_components=2)),
                        ('nb_classifier', GaussianNB())])
pipeline_nb.fit(X_train, y_train)
score4 = pipeline_nb.score(X_test, y_test)
pred4 = pipeline_nb.predict(X_test)
acc4 = metrics.accuracy_score(y_test, pred4)
print('Classification report:')
print(metrics.classification_report(y_test, pred4))
print('Accuracy score: ', acc4)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/naive_bayes.py:206: DataConversionWarning: A column-ve
ctor y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for ex
ample using ravel().
  y = column_or_1d(y, warn=True)
Classification report:
precision    recall  f1-score   support
1         0.00         0.00         0.00         59804
2         0.50         1.00         0.67         179704
3         0.00         0.00         0.00         120492
accuracy          0.25          0.33          0.33         360000
macro avg         0.17         0.33         0.22         360000
weighted avg         0.25         0.50         0.33         360000
Accuracy score: 0.4991777777777778
```

We note that changing the model gives us very similar results. This is probably due to the fact that there were no evident correlations to learn between the features and the target variable and hence changing the type of learning function or the shape of the line of fit would not necessarily improve the model.

## Treating the multi-class classification problem as several binary classification problems

The next thing I tried was to split the multi-class classification procedure into steps of binary classifications. Implying to train a model to learn to classify instances to 'Class A' and 'Not Class A' instead of 'Class A', 'Class B', and 'Class C'. Similarly for 'Class B'/'Class C' and 'Not Class B'/'Class C'.

Model 1: [1, not 1]  
Model 2: [2, not 2]  
Model 3: [3, not 3]

Let us first try to build a model that can classify instances into 'Class 1' and 'not Class 1'

```
In [ ]: df_copy1 = df.copy('Class1' == 2) | (df_copy1['Class1' == 3], 'Class' ) = 0
```

```
In [ ]: df_copy1.head()
```

```
Out [ ]:
   A      B      C      D      E      F      G      H      I      J      K
0 231.420023 217.624839 -15.611916 198.160805 82.873279 127.350084 224.592828 -5
1 -38.019270 -14.185695 9.583547 22.293822 -25.578283 -18.373955 -0.094457 -33.711852 -8.356041 23.782402 4.199023 2
2 -39.197085 -20.418850 21.023083 19.790280 -25.902587 -19.189004 -2.953836 -25.299219 -6.612401 26.285392 5.911292 6
3 221.630408 -5.785352 21.672532 -9.900781 126.795177 85.122288 106.857593 197.640135 82.560019 157.105143 212.889231 -3
4 228.558412 -12.447710 204.637218 -13.277704 138.930529 91.101870 115.588954 209.300011 89.961688 130.299732 201.795100 -1
```

```
In [ ]: df_1 = df_copy1
```

```
In [ ]: df_1['Class1'].unique()
```

```
Out [ ]: array([0, 1])
```

```
In [ ]: X1 = df_1.drop(['Class1'], axis=1)
y1 = df_1['Class1']
```

```
In [ ]: X_train1, X_test1, y_train1, y_test1 = train_test_split(X1, y1, stratify=y, test_size=0.2)
pipeline_rf1 = Pipeline([('scaler', StandardScaler()),
                          ('pca', PCA(n_components=2)),
                          ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf1.fit(X_train1, y_train1)
pipeline_rf1.score(X_test1, y_test1)
```

```
Out [ ]: 0.8334166666666666
```

We see a better performance in the model pipeline when we train it on a single target class by considering the other two classes as a single entity in this case, '1 and not 1'.

Let us now try and model for classes 2 and 3.

```
In [ ]: df_original.head()
```

```
Out [ ]:
   A      B      C      D      E      F      G      H      I      J      K
0 231.420023 217.624839 -15.611916 198.160805 82.873279 127.350084 224.592828 -5
1 -38.019270 -14.185695 9.583547 22.293822 -25.578283 -18.373955 -0.094457 -33.711852 -8.356041 23.782402 4.199023 2
2 -39.197085 -20.418850 21.023083 19.790280 -25.902587 -19.189004 -2.953836 -25.299219 -6.612401 26.285392 5.911292 6
3 221.630408 -5.785352 21.672532 -9.900781 126.795177 85.122288 106.857593 197.640135 82.560019 157.105143 212.889231 -3
4 228.558412 -12.447710 204.637218 -13.277704 138.930529 91.101870 115.588954 209.300011 89.961688 130.299732 201.795100 -1
```

```
In [ ]: df_2 = df_copy1
```

```
In [ ]: df_2.loc[(df_2['Class1'] == 1) | (df_2['Class1'] == 3), 'Class' ] = 0
```

```
In [ ]: df_2.head()
```

```
Out [ ]:
   A      B      C      D      E      F      G      H      I      J      K
0 231.420023 217.624839 -15.611916 198.160805 82.873279 127.350084 224.592828 -5
1 -38.019270 -14.185695 9.583547 22.293822 -25.578283 -18.373955 -0.094457 -33.711852 -8.356041 23.782402 4.199023 2
2 -39.197085 -20.418850 21.023083 19.790280 -25.902587 -19.189004 -2.953836 -25.299219 -6.612401 26.285392 5.911292 6
3 221.630408 -5.785352 21.672532 -9.900781 126.795177 85.122288 106.857593 197.640135 82.560019 157.105143 212.889231 -3
4 228.558412 -12.447710 204.637218 -13.277704 138.930529 91.101870 115.588954 209.300011 89.961688 130.299732 201.795100 -1
```

```
In [ ]: df_2['Class'].unique()
```

```
Out [ ]: array([0, 1])
```

```
In [ ]: X2 = df_2.drop(['Class1'], axis=1)
y2 = df_2['Class1']
```

```
In [ ]: X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2, stratify=y, test_size=0.2)
pipeline_rf2 = Pipeline([('scaler', StandardScaler()),
                          ('pca', PCA(n_components=2)),
                          ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf2.fit(X_train2, y_train2)
pipeline_rf2.score(X_test2, y_test2)
```

```
Out [ ]: 0.499975
```

```
In [ ]: df_3 = df_original.copy()
df_3.loc[(df_3['Class1'] == 1) | (df_3['Class1'] == 2), 'Class' ] = 0
df_3['Class1'].unique()
```

```
Out [ ]: array([0, 3])
```

```
In [ ]: X3 = df_3.drop(['Class1'], axis=1)
y3 = df_3['Class1']
```

```
In [ ]: X_train3, X_test3, y_train3, y_test3 = train_test_split(X3, y3, stratify=y, test_size=0.2)
pipeline_rf3 = Pipeline([('scaler', StandardScaler()),
                          ('pca', PCA(n_components=2)),
                          ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf3.fit(X_train3, y_train3)
pipeline_rf3.score(X_test3, y_test3)
```

```
Out [ ]: 0.6660166666666667
```

As we can see, training binary classifiers on our data seems to improve the accuracy to a certain extent.

Model accuracies:

- 1 and not 1: 83.33%
- 2 and not 2: 49.99%
- 3 and not 3: 66.6%

Let us now combine these pipelines into a single model pipeline that can hopefully fit the entire dataset. We use sklearn's VotingClassifier to combine multiple models.

```
In [ ]: from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('rf1', pipeline_rf1), ('rf2', pipeline_rf2), ('rf3', pipeline_rf
3)], voting='hard')
```

```
In [ ]: vclf.fit(X_train, y_train)
vclf.score(X_test, y_test)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:235: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:268: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
Out [ ]: 0.4991777777777778
```

```
In [ ]: pipeline_vclf = Pipeline([('scaler', StandardScaler()),
                                ('pca', PCA(n_components=2)),
                                ('vclf_classifier', VotingClassifier(estimators=[('rf1', pipeline_rf1), ('rf2', pip
eline_rf2), ('rf3', pipeline_rf3)], voting='hard'))])
pipeline_vclf.fit(X_train, y_train)
vclf.score(X_test, y_test)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:235: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:268: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
Out [ ]: 0.4991777777777778
```

Now, my hypothesis was that considering the fact that most of the instances in our main dataset are classified into class 2, we could create two binary classification models: one that classifies instances for just target 'class 2' and 'not class 2' and another model that classifies 'class 1' and 'class 3'. And then combining these two models to make predictions on our main dataset could help create a working model.

```
In [ ]: ## Model to classify 1 and 3
df_13 = df_original.copy()
df_13.drop(df_13[df_13['Class1'] == 2].index, inplace=True)
df_13['Class1'].unique()
```

```
Out [ ]: array([3, 1])
```

```
In [ ]: len(df_13)
```

```
Out [ ]: 600772
```

```
In [ ]: X13 = df_13.drop(['Class1'], axis=1)
y13 = df_13['Class1']
```

```
In [ ]: X_train13, X_test13, y_train13, y_test13 = train_test_split(X13, y13, test_size=0.2)
```

```
In [ ]: X_train13, X_test13, y_train13, y_test13 = train_test_split(X13, y13, test_size=0.2)
pipeline_rf13 = Pipeline([('scaler', StandardScaler()),
                           ('pca', PCA(n_components=2)),
                           ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf13.fit(X_train13, y_train13)
pipeline_rf13.score(X_test13, y_test13)
```

```
Out [ ]: 0.686613124713912
```

Now, we combine this model with the one that was able to classify class 2 alone and the model the train dataset.

```
In [ ]: pipeline_vclf = Pipeline([('scaler', StandardScaler()),
                                ('pca', PCA(n_components=2)),
                                ('vclf_classifier', VotingClassifier(estimators=[('rf13', pipeline_rf13), ('rf2', p
ipeline_rf2)], voting='hard'))])
pipeline_vclf.fit(X_train, y_train)
vclf.score(X_test, y_test)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:235: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:268: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
Out [ ]: 0.4991777777777778
```

We notice the same accuracy as before.

Finally, I tried to combine the models that can classify [Class 2, Class 3], [Class 2, Class 1] and [Class 1 and Class 3]

```
In [ ]: ## Class 1 and Class 2
df_12 = df_original.copy()
df_12.drop(df_12[df_12['Class1'] == 3].index, inplace=True)
df_12['Class1'].unique()
```

```
Out [ ]: array([2, 1])
```

```
In [ ]: X12 = df_12.drop(['Class1'], axis=1)
y12 = df_12['Class1']
```

```
In [ ]: X_train12, X_test12, y_train12, y_test12 = train_test_split(X12, y12, test_size=0.2)
```

```
In [ ]: X_train12, X_test12, y_train12, y_test12 = train_test_split(X12, y12, test_size=0.2)
pipeline_rf12 = Pipeline([('scaler', StandardScaler()),
                           ('pca', PCA(n_components=2)),
                           ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf12.fit(X_train12, y_train12)
pipeline_rf12.score(X_test12, y_test12)
```

```
Out [ ]: 0.7499499512024224
```

```
In [ ]: ## Class 2 and 3
df_23 = df_original.copy()
df_23.drop(df_23[df_23['Class1'] == 1].index, inplace=True)
df_23['Class1'].unique()
```

```
Out [ ]: array([2, 3])
```

```
In [ ]: X23 = df_23.drop(['Class1'], axis=1)
y23 = df_23['Class1']
```

```
In [ ]: X_train23, X_test23, y_train23, y_test23 = train_test_split(X23, y23, test_size=0.2)
```

```
In [ ]: X_train23, X_test23, y_train23, y_test23 = train_test_split(X23, y23, test_size=0.2)
pipeline_rf23 = Pipeline([('scaler', StandardScaler()),
                           ('pca', PCA(n_components=2)),
                           ('rf_classifier', RandomForestClassifier(max_depth=5))])
pipeline_rf23.fit(X_train23, y_train23)
pipeline_rf23.score(X_test23, y_test23)
```

```
Out [ ]: 0.5996990030939699
```

Now we combine the three models.

## Final Model

```
In [ ]: pipeline_vclf123 = Pipeline([('scaler', StandardScaler()),
                                    ('pca', PCA(n_components=2)),
                                    ('vclf_classifier', VotingClassifier(estimators=[('rf13', pipeline_rf13), ('rf12', p
ipeline_rf12), ('rf23', pipeline_rf23)], voting='hard'))])
pipeline_vclf123.fit(X_train, y_train)
vclf123.score(X_test, y_test)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:235: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.7/dist-packages/sklearn/preprocessing/_label.py:268: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,
), for example using ravel().
  y = column_or_1d(y, warn=True)
Out [ ]: 0.4991777777777778
```

This is the final model that we will use to make predictions.

## ONNX

```
In [ ]: from sklearn import convert_sklearn
from sklearn.common.data_types import FloatTensorType
onnx_model_path = '/content/drive/MyDrive/Datasets/AkhilSrinath_final_model.onnx'
```

```
num_features = 15
initial_type = [('float_input', FloatTensorType([None, num_features]))]
onnx = convert_sklearn(pipeline_vclf123, initial_types=initial_type)
with open(onnx_model_path, 'wb') as f:
    f.write(onnx.SerializeToString())
```

```
RuntimeError                                Traceback (most recent call last)
/usr/local/lib/python3.7/dist-packages/sklearn/_parse.py in parse_sklearn(scope, model, inputs, cust
om_parsers, final_types)
--> 519         output = parse_sklearn(
520             scope.temp(), model, inputs, custom_parsers=custom_parsers)
521         except RuntimeError:
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/_parse.py in _parse_sklearn(scope, model, inputs, cus
tom_parsers, final_types)
466         outputs = sklearn_parsers_map[model](scope, model, inputs,
467         custom_parsers=custom_parsers)
--> 488         elif isinstance(model, pipeline.Pipeline):
489             parser = sklearn_parsers_map[pipeline.Pipeline]
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/_parse.py in _parse_sklearn_pipeline(scope, model, in
puts, custom_parsers)
195         inputs = _parse_sklearn(scope, step[1], inputs,
196         custom_parsers=custom_parsers)
--> 197         return inputs
198         custom_parsers=custom_parsers)
199         return inputs
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/_parse.py in _parse_sklearn(scope, model, inputs, cus
tom_parsers, final_types)
466         outputs = sklearn_parsers_map[model](scope, model, inputs,
467         custom_parsers=custom_parsers)
--> 488         elif isinstance(model, pipeline.Pipeline):
489             parser = sklearn_parsers_map[pipeline.Pipeline]
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/_parse.py in _parse_sklearn_pipeline(scope, model, in
puts, custom_parsers)
195         inputs = _parse_sklearn(scope, step[1], inputs,
196         custom_parsers=custom_parsers)
--> 197         return inputs
198         custom_parsers=custom_parsers)
199         return inputs
```

```
/usr/local/lib/python3
```